

Analyzing Hate Crime Trends in the United States

Ria Mathew, Faith Ang, Xuanwei Ding

November 7, 2025

##“Which groups (racial, religious, gender, state etc. . .) are most frequently targeted, and what may have caused the increase or decrease of certain types of hate crime over the years?”

Abstract

The purpose of this research and statistical analysis is to begin identifying significant trends in hate crime incidents across the United States, focusing on bias motivation, incident year, location and offense type. Using cleaned, standardized hate crime data, we analyze how these incidents have changed over time and which categories of hate crimes are most prevalent and where these cases appear the most frequently. Our data indicates that certain biases are more targeted (racial), and even within each bias (race/ethnicity, religion, sexual orientation, other) subgroups are also more targetted than others for instance Anti-Gay sentiment within the sexual orientation bias. Although the number of hate crimes fluctuates over time, there are several categories with clear trends that may reflect coexisting social and political dynamics going on at the same time.

Using data visualizations such as bar graphs and statistical summaries (such as tables), we can identify hate crime trends and draw comparisons by bias, offender’s race, and other variables. Our hope is to provide valuable insights that will lead us closer to discovering how bias-related crimes have evolved over time, discover what may trigger these crimes, and see who may be most susceptible to be victim of such hate crimes. We hope to find how data and statistical analysis can be used to make social justice discoveries.

Introduction

The purpose of this analysis is to investigate patterns in hate crimes across the United States. To answer our overarching research question, we developed specific prompts to guide our analysis and the visualizations we will create. These 6 prompts are also written out and further explored in our Plots.Rmd file.

1. Count the number of offenses per bias to see which group is most frequently targeted
2. Compare the number of offenses by race to identify which groups are more likely to be the victim to what offenses.
3. Analyze how the number of offenses has changed over time.
4. Examine variation by state of offense types
5. See how many of each type of offense there is and compare quantities
6. Find correlations between biases and offense (eg. do Asians experience intimidation based hate crimes or assault based hate crimes)

The results of this analysis may benefit policymakers, sociologists, and law enforcement agencies and even the FBI from which the data set itself was retrieved. It may benefit them by identifying demographic trends and creating strategies towards prevention. Also understanding which groups may be most at risk can even help larger organizations and employers to create resources for employees to easily report such incidents.

Data

The data we used was the `hate_crime.csv` dataset. The raw data was an FBI dataset (retrieved directly from their website) which contains data over the span of 33 years (1994 to 2024) with 265,834 individual hate crime cases with 28 variables ranging from the dates of the incidents, the perpetrators race, the crime location and more. For this study we focused on the data that summarized hate crime cases by year and bias category, disregarding agency information and age-related data.

Data Cleaning

The data cleaning process had a couple steps that helped ensure the data was consistent, generalization, relevant, and filled. First, any rows with missing values were removed. Next, all character-type columns were converted to factors for statistical analysis and visualization of categorical data such as the variables race, offense type, and bias. Finally, several columns were renamed for readability. These steps ensured that the dataset was clean, consistent, and ready for analysis.

After cleaning we decided to focus on the variables: - Race of Offender: race of the offender

- Year: year of the offense


- State: where the offense took place

- Bias: the group targeted by the hate crime

- Offense: type of crime (e.g., intimidation, assault)

- Victim Count: number of victims per case


Visualizations and Analysis



To begin our analysis, it was very important that we knew which group was the most frequently the victim of hate crime. We did this by categorizing the Bias data into 4 primary categories: Race/Ethnicity, Religion, Sexual Orientation, Other. Other consisted of biases that did not fall into these prominent categories yet

did not have consistent data to create a category of its own, such as gender-affiliated ones. Once separating based on bias, we can see the incident count for racially motivated hate crime is almost 3.5 times as often as religion and sexual orientation which have roughly the same amount of reported incidences.

We conclude that a majority of hate crime targets race. This is not unexpected, as Race/Ethnicity are often more immediately visible attributes, whereas other identity characteristics such as religion and sexual orientation may not be as obvious.



To better understand which groups are most targeted within each broad bias category, we filtered the dataset to include only incidents related to these 3 main groups: Race/Ethnicity, Religion, and Sexual Orientation. For each category, broke down the data even further by specific subBias groups (ex: Anti-Black, Anti-Jewish, Anti-Gay) and plotted the incident count for each subgroup. By breaking these down we aim to find which groups are most targeted.

As a side note, through the rest of this report we will be addressing more encompassing biases as ‘main bias’ groups (Race/Ethnicity, Religion, and Sexual Orientation) and the sub categories within these 3 aforementioned bias groups as ‘sub bias’ groups (ex: Anti-Black within Race/Ethnicity or Anti-Jewish within Religion).


Across the 3 bar graphs we see the pattern that one group has a large portion of the reported hate crimes, showing strong disproportions between each sub bias of our 3 main bias categories.

Race/Ethnicity Main Bias: This plot shows highlights that there may be Anti-Black or African American sentiment since this demographic represents the largest share of hate crime victims. Approximating 88,000 reported incidents disproportionate to other racial and ethnic groups which range from a few thousand to just over 25,000.

Religion Main Bias: Considering ratios, it seems an overwhelming amount of attacks are towards Jewish individuals. It may also be important to note that the identity ‘Jewish’ could also pertain to someone’s ethnicity so this classification may contain an even wider range of incidents.

Sexual Orientation Main Bias: We see the most affected group is Gay (male) individuals, followed by a broad category that contains Lesbians, Gay, Bisexual, or Transgender individuals, what is clarified as a mixed group, less clearly defined since it broadly contains many subgroups.


Each of our main bias categories has one sub-bias group that clearly stands out as the most targeted. For Race/Ethnicity it was Anti-Black or African American, for religion it was Anti-Jewish, for sexual orientation it was Gay (Male). Across all three plots, we see a concentrated pattern of hate crimes disproportionately affecting a single community within each category, highlighting substantial disparities between sub bias groups. While our earlier graph showed that Race/Ethnicity has a significantly higher overall number of incidents between our main bias groups, this more detailed breakdown reveals how hate crimes are similarly not evenly distributed within each main bias. The outlier subgroups that we see in each of our main biases we plotted displays an uneven distribution in levels of victimization.



Looking further into Race/Ethnicity, since it has the largest number of reported incidents by category, we broke down each group by the proportion of different hate-crime types. To make sure that our comparisons were not skewed by differences in total incident counts between groups, we normalized the data by scaling it so each bar reflects the proportion instead of raw incident count. Among all hate-crime categories, the five reoccurring ones we identified were Intimidation, Destruction/Damage/Vandalism of Property, Simple Assault, Aggravated Assault, and Burglary/Breaking & Entering.

With each of the different races having similar offense type breakdowns (similar proportions of each offense type), the normalized graph typically assists in highlighting differences. However, from the little variations we see, we may be able to conclude that offense types may not be specific to race as the differences are miniscule and may just be chance and due to smaller sample sizes in some of the races.

From this graph, the graph with the least similar proportions are that from “Anti-Multiple Races, Group” which we can’t accurately say reflects differences in offense type based on race since it includes multiple races.




These visuals track hate crime incidents over the years, with the second displaying the three most targeted sub-bias groups from our previously data visuals. Each point represents a year and their size is corresponds with the number of reported incidents.


The first line plot shows the overall trend with no bias specific statistics showing upswings 2018 onwards and similarly more incidents between 1996 and 2001. It also shows a significant dip between 2008 and 2017. With our second bubble plot we can compare and see if this trend is reflected among the subgroups that most dominate our data set: our 3 prominent sub biases from each main category (Anti—Black or African American, Anti-Jewish, Anti-Gay.

Although it is a little difficult to tell since there are more incidents with our Race main Bias in comparison with our other bias categories, we can see that the bubble graph still reflects the same general trends. There are larger bubbles across all three sub biases during the same years as those reflected in the line graph, and visa versa with the smaller bubbles. This is most visible with our Anti—Black or African American sub bias group, but that may just be more pronounced because there are more incidents while the other 2 sub bias groups are more subtle.

Based on the time charts, we can infer that overall hate-crime incidents are influenced by social or historical forces that affect all groups, not just one. For instance, the upswing from 2018 onwards may be due to the tensions with the political climate or increased media presence, or it may be as simple as better methods for reporting incidents where it seems as though hate crime is increasing when in reality it is changes in accessibility.



Hate Crime Incidents Over Time



To deepen our understanding of hate crime patterns, we aimed to see how location shapes hate crimes. Examining location is critical because hate crimes are not randomly distributed across the US. They are may be more concentrated in some states because of differences in population size, demographic diversity, political environments etc. Identifying these geographic patterns helps policymakers and communities create

prevention resources based on where they are most needed.

The West and South consistently report the highest case rates mostly because they have the most populous states (California, Texas, Florida, etc.). The Northeast and Midwest have lower totals, which could be due to having fewer states or smaller and less diverse areas. Despite these differences, all four regions have similar long-term trends: a peak in the late 1990s to early 2000s, a decline from 2008–2016, and an increase from around 2017–2018 to 2024. These patterns can help us determine how national level events such as political polarization, incidents that are popularized by social media, leadership changes, or, as mentioned before, simply improving the reporting mechanisms can have a powerful influence across the nation.

California, New York, New Jersey, Washington, and Massachusetts appear near the top in nearly every category which may reflect their large populations and maybe more established hate-crime reporting infrastructure. We also see that state by state there are differences even between the three categories. For instance, New Jersey has high intimidation numbers but much lower simple-assault counts. On the other hand, there are some states that show little relative variation across all offense types like Texas and Florida which indicate a more balanced distribution of hate-crime expressions.

Chi-Squared Test

Pearson's Chi-squared test

data: table_pv X-squared = 8326.1, df = 24, p-value < 2.2e-16

This chi-square test examined whether the distribution of the four most common offense types (vandalism, intimidation, simple assault, and aggravated assault) differed across victim bias categories. From this we got that there is a statistically significant relationship between the victim's targeted group and the type of offense committed since $p\text{-value} < 2.2e-16$ which is far smaller than our significance level $p = 0.05$.

This tells us that the patterns for offenses are not random when considering different biases. Certain types of hate-motivated offenses occur more frequently against particular more targeted groups. Some victim groups may experience more intimidation incidents, while others face higher levels of physical assaults. The results suggest that these differences are meaningful and can be used to make policy decisions to protect the more vulnerable communities.

Conclusion

The purpose of this research was to find hate crime patterns, examine deeper the bias motivation, offense type, location, and how things change over time. After conducting our analysis, it revealed that certain biases, especially ethnic motivations, were the largest share of hate crimes, consequently telling us they were most targeting. Within broader main biases, specific sub bias groups were disproportionately targeted like such as Anti-Gay (male), Anti-Black and African American and Anti-Jewish. We also observed that incident counts fluctuating across the years which may be due to broad social or political events. Through data visualizations and our chi-squared test, we were able to compare trends across states and biases and highlight how hate crime patterns change across demographically and geographically.

For continued research, we can see how social, economic, or political variables determine what triggers increases or decreases in hate crimes. We can also see what causes under reporting/increased reporting in different locations. All in all, this project shows the value of using statistics to uncover meaningful social justice trends to support policymakers and protect the community as a whole.

Contributions Section

THE GitHub link:

{r} Project

Faith Ang: - Created the GitHub repository and the R and Rmd blank files for the project - Wrote a substantial amount of the README.md file - Made ggplot templates for various plots in 02_funcut_Plots.R, and functions for data cleaning - Worked on 11_DataCleaning.Rmd and 12_Plots.Rmd (question 1). - Worked on the Visualization and Analysis portion of FinalReport.Rmd

Ria Mathew: - Edited: README.md, 01_funcut_DataCleaning.R, 02_funcut_Plots.R. (grammar, spelling, and smoother transitions) - Worked on 11_DataCleaning.Rmd (revised and improved) and 12_Plots.Rmd (question 2). - Worked on Abstract, Introduction, Data, Data Cleaning, Visualization and Analysis, Chi-Squared Test, Conclusion, portion of FinalReport.Rmd.

Xuanwei Ding: - Worked substantially on 12_Plots.Rmd, coding for questions 3, 4, 5, and 6 on 12_Plots.Rmd. - Worked on the analysis of the top 3 hate crimes.